

```
In [1]: #data preperation and analysis library
import pandas as pd

#plotting libraries
import matplotlib.pyplot as plt
import seaborn as sns

#library for creating random samples
from sklearn.model_selection import train_test_split

#library for buliding linear regression model
from sklearn.linear_model import LinearRegression

#feature selection (to select significant variables)
from sklearn.feature_selection import SelectKBest, f_regression
```

```
In [2]: #Load data
df=pd.read_csv(r"C:\Users\ACER\Desktop\introtallent\python\data\104380_Python
```

```
In [3]: df.head(2)
```

Out[3]:

	MPG	Cylinders	Displacement	Horsepower	Weight	Acceleration	Model_year	Origin	Car_N
0	8.0	8	307.0	130	3504	12.0	2015	1	chev
1	15.0	8	350.0	165	3693	11.5	2015	1	t

```
In [4]: df=df.drop("Car_Name",axis=1)
```

```
In [5]: df=df.drop("Model_year",axis=1)
```

```
In [6]: df.head()
```

Out[6]:

	MPG	Cylinders	Displacement	Horsepower	Weight	Acceleration	Origin
0	8.0	8	307.0	130	3504	12.0	1
1	15.0	8	350.0	165	3693	11.5	1
2	18.0	8	318.0	150	3436	11.0	1
3	16.0	8	304.0	150	3433	12.0	1
4	17.0	8	302.0	140	3449	10.5	1

```
In [7]: #print row and column count  
df.shape
```

```
Out[7]: (398, 7)
```

```
In [8]: df.dtypes
```

```
Out[8]: MPG                float64  
Cylinders                int64  
Displacement            float64  
Horsepower              object  
Weight                  int64  
Acceleration            float64  
Origin                  int64  
dtype: object
```

```
In [9]: #cylinders,Origin are categorical variables stored as int  
#change the datatype to object  
df['Cylinders']=df['Cylinders'].astype('object')  
df['Origin']=df['Origin'].astype('object')
```

```
In [10]: df['Horsepower']=pd.to_numeric(df['Horsepower'],errors='coerce')
```

```
In [11]: df.dtypes
```

```
Out[11]: MPG                float64  
Cylinders                object  
Displacement            float64  
Horsepower              float64  
Weight                  int64  
Acceleration            float64  
Origin                  object  
dtype: object
```

```
In [12]: #Feature engineering-[check and impute missing values,if any]  
df.isnull().sum()
```


```
Out[12]: MPG                0  
Cylinders                0  
Displacement            0  
Horsepower              6  
Weight                  0  
Acceleration            0  
Origin                  0  
dtype: int64
```

```
In [13]: #Impute Horsepower with median  
df['Horsepower']=df['Horsepower'].fillna(df['Horsepower'].median())
```


```
In [14]: df.isnull().sum()
```

```
Out[14]: MPG          0  
Cylinders          0  
Displacement       0  
Horsepower         0  
Weight             0  
Acceleration       0  
Origin             0  
dtype: int64
```


```
In [15]: #feature engineering-outlier treatment  
plt.boxplot(df['MPG'])  
plt.show()
```




```
In [16]: plt.boxplot(df['Cylinders'])  
plt.show()
```




```
In [17]: plt.boxplot(df['Displacement'])  
plt.show()
```



```
In [18]: plt.boxplot(df['Weight'])  
plt.show()
```




```
In [19]: plt.boxplot(df['Acceleration'])  
plt.show()
```




```
In [20]: #user defined function to remove outliers
def remove_outliers(d,c):
    q1=d[c].quantile(0.25)
    q3=d[c].quantile(0.75)
    iqr=q3-q1
    ub=q3+1.5*iqr
    lb=q1-1.5*iqr
    #remove outliers and store good data in result
    result=d[(d[c]>=lb) & (d[c]<=ub)]
    return result
```

```
In [22]: #remove outliers from Acceleration
df=remove_outliers(df,'Acceleration')
plt.boxplot(df['Acceleration'])
plt.show()
```



```
In [25]: #remove outliers from Horsepower
df=remove_outliers(df,'Horsepower')
plt.boxplot(df['Horsepower'])
plt.show()
```



EDA(Exploratory Data Analysis)


```
In [26]: #create pairplot
sns.pairplot(df)
```

```
Out[26]: <seaborn.axisgrid.PairGrid at 0x1683cf5f7f0>
```




```
In [27]: #data mix
# 'Cylinders', 'Origin'
```

```
In [28]: df.groupby('Cylinders')['Cylinders'].count().plot(kind='bar')  
plt.show()
```




```
In [29]: df.groupby('Origin')['Origin'].count().plot(kind="bar")  
plt.show()
```



```
In [30]: #distribution
sns.distplot(df['MPG'])
plt.show()
```

C:\Users\ACER\anaconda3\lib\site-packages\seaborn\distributions.py:2551: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


warnings.warn(msg, FutureWarning)



```
In [31]: sns.distplot(df['Displacement'])
plt.show()
```

C:\Users\ACER\anaconda3\lib\site-packages\seaborn\distributions.py:2551: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


warnings.warn(msg, FutureWarning)




```
In [32]: sns.distplot(df['Weight'])  
plt.show()
```

C:\Users\ACER\anaconda3\lib\site-packages\seaborn\distributions.py:2551: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


warnings.warn(msg, FutureWarning)



```
In [33]: sns.distplot(df['Acceleration'])  
plt.show()
```

C:\Users\ACER\anaconda3\lib\site-packages\seaborn\distributions.py:2551: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


warnings.warn(msg, FutureWarning)



```
In [34]: sns.distplot(df['Horsepower'])
plt.show()
```

C:\Users\ACER\anaconda3\lib\site-packages\seaborn\distributions.py:2551: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```



```
In [35]: #check object variables for spelling differences, and redundant data
#Cylinders object
#Origin object
```

```
In [36]: df['Cylinders'].unique()
```

```
Out[36]: array([8, 4, 6, 3, 5], dtype=object)
```

```
In [37]: df['Origin'].unique()
```

```
Out[37]: array([1, 3, 2], dtype=object)
```

Feature Engineering:One-hot-encoding(dummy conversion)

```
In [38]: #store all categorical variables in a new dataframe
df_categorical=df.select_dtypes(include=['object'])
```

```
In [39]: df_categorical.head()
```

Out[39]:

	Cylinders	Origin
0	8	1
2	8	1
3	8	1
4	8	1
12	8	1

```
In [40]: #create dummy
dummy=pd.get_dummies(df_categorical,drop_first=True)
```

```
In [41]: dummy.head()
```

Out[41]:

	Cylinders_4	Cylinders_5	Cylinders_6	Cylinders_8	Origin_2	Origin_3
0	0	0	0	1	0	0
2	0	0	0	1	0	0
3	0	0	0	1	0	0
4	0	0	0	1	0	0
12	0	0	0	1	0	0

```
In [42]: #combine numeric columns from df with dummy columns to create final data
df_numeric=df.select_dtypes(include=['int64','float64'])
```

```
In [43]: df_master=pd.concat([df_numeric,dummy],axis=1)
```

```
In [44]: df_master.head()
```

Out[44]:

	MPG	Displacement	Horsepower	Weight	Acceleration	Cylinders_4	Cylinders_5	Cylinders_6	Cylinders_8
0	8.0	307.0	130.0	3504	12.0	0	0	0	1
2	18.0	318.0	150.0	3436	11.0	0	0	0	1
3	16.0	304.0	150.0	3433	12.0	0	0	0	1
4	17.0	302.0	140.0	3449	10.5	0	0	0	1
12	15.0	400.0	150.0	3761	9.5	0	0	0	1



**Create X (with all independent variables) and Y
(With the target variable)**

```
In [45]: x=df_master.drop('MPG',axis=1)
```

```
In [46]: y=df_master['MPG']
```

Random sampling: create training and test samples

```
In [47]: #create training and test samples  
xtrain,xtest,ytrain,ytest=train_test_split(x,y,train_size=0.7,random_state=0)
```

Feature Selection

```
In [48]: #create key_features object to select the top k features  
# key_features = SelectKBest(score_func=f_regression,k='all')  
key_features=SelectKBest(score_func=f_regression,k=5)  
#Fit the key_features to the training data and transform it  
xtrain_selected=key_features.fit_transform(xtrain,ytrain)  
#Get the indices of the selected features  
selected_indices=key_features.get_support(indices=True)  
#Get the names of the selected features  
selected_features=xtrain.columns[selected_indices]
```

```
In [49]: selected_features
```

```
Out[49]: Index(['Displacement', 'Horsepower', 'Weight', 'Cylinders_4', 'Cylinders_8'],  
dtype='object')
```

Instantiate linear regression

```
In [50]: linreg=LinearRegression()
```

Model 1: Build training model using all features

```
In [51]: linreg.fit(xtrain,ytrain)
```

```
Out[51]: LinearRegression()
```

```
In [52]: linreg.score(xtrain,ytrain)
```

```
Out[52]: 0.731954532119161
```

```
In [53]: #test yours model's Learning
#predict mileage using test samples
predicted_mileage=linreg.predict(xtest)
```

```
In [54]: #check r-squared (accuracy score)
linreg.score(xtest,ytest)
```

Out[54]: 0.7329949257086605

Model 2: Build model using KBest selected feaures

```
In [55]: ##store KBest columns from xtrain to xtrain_kbest
xtrain_kbest=xtrain[selected_features]
```

```
In [56]: xtrain_kbest.head()
```

Out[56]:

	Displacement	Horsepower	Weight	Cylinders_4	Cylinders_8
334	70.0	100.0	2420	0	0
175	90.0	70.0	1937	1	0
112	122.0	85.0	2310	1	0
2	318.0	150.0	3436	0	1
198	91.0	53.0	1795	1	0

```
In [57]: linreg.fit(xtrain_kbest,ytrain)
```

Out[57]: LinearRegression()

```
In [58]: linreg.score(xtrain_kbest,ytrain)
```

Out[58]: 0.7110901050565608

```
In [59]: #store KBest columns from xtest to xtest_kbest
xtest_kbest=xtest[selected_features]
```

```
In [60]: pred_y=linreg.predict(xtest_kbest)
```

```
In [61]: linreg.score(xtest_kbest,ytest)
```

Out[61]: 0.6951570744108746


log transformation

BUILT MODEL AFTER REDUCING SKEWNESS IN THE DISPLACEMENT AND WEIGHT VARIABLE

```
In [63]: sns.distplot(df['Displacement'])
plt.show()
```

C:\Users\ACER\anaconda3\lib\site-packages\seaborn\distributions.py:2551: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```



```
In [64]: import numpy as np
df['log_Displacement']=np.log(df['Displacement'])
```

```
In [65]: df.head(3)
```

Out[65]:


	MPG	Cylinders	Displacement	Horsepower	Weight	Acceleration	Origin	log_Displacement
0	8.0	8	307.0	130.0	3504	12.0	1	5.726848
2	18.0	8	318.0	150.0	3436	11.0	1	5.762051
3	16.0	8	304.0	150.0	3433	12.0	1	5.717028

```
In [66]: sns.distplot(df['log_Displacement'])
```

C:\Users\ACER\anaconda3\lib\site-packages\seaborn\distributions.py:2551: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```


```
Out[66]: <AxesSubplot:xlabel='log_Displacement', ylabel='Density'>
```



```
In [67]: sns.distplot(df['Weight'])  
plt.show()
```

C:\Users\ACER\anaconda3\lib\site-packages\seaborn\distributions.py:2551: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```



```
In [69]: import numpy as np
df['log_Weight']=np.log(df['Weight'])
```

```
In [71]: df.head(3)
```

Out[71]:


	MPG	Cylinders	Displacement	Horsepower	Weight	Acceleration	Origin	log_Displacement
0	8.0	8	307.0	130.0	3504	12.0	1	5.726848
2	18.0	8	318.0	150.0	3436	11.0	1	5.762051
3	16.0	8	304.0	150.0	3433	12.0	1	5.717028

```
In [72]: sns.distplot(df['log_Weight'])
```

C:\Users\ACER\anaconda3\lib\site-packages\seaborn\distributions.py:2551: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)

Out[72]: <AxesSubplot:xlabel='log_Weight', ylabel='Density'>



```
In [73]: df=df.drop(['Displacement','Weight'],axis=1)
df.head()
```

Out[73]:

	MPG	Cylinders	Horsepower	Acceleration	Origin	log_Displacement	log_Weight
0	8.0	8	130.0	12.0	1	5.726848	8.161660
2	18.0	8	150.0	11.0	1	5.762051	8.142063
3	16.0	8	150.0	12.0	1	5.717028	8.141190
4	17.0	8	140.0	10.5	1	5.710427	8.145840
12	15.0	8	150.0	9.5	1	5.991465	8.232440

In [74]: `df.dtypes`

Out[74]:

MPG	float64
Cylinders	object
Horsepower	float64
Acceleration	float64
Origin	object
log_Displacement	float64
log_Weight	float64
dtype:	object

In [75]: `dummy.head()`

Out[75]:

	Cylinders_4	Cylinders_5	Cylinders_6	Cylinders_8	Origin_2	Origin_3
0	0	0	0	1	0	0
2	0	0	0	1	0	0
3	0	0	0	1	0	0
4	0	0	0	1	0	0
12	0	0	0	1	0	0

In [76]: *#combine numeric columns from df with dummy columns to create final data*
`df_numeric=df.select_dtypes(include=['int64','float64'])`

In [77]: `df_master=pd.concat([df_numeric,dummy],axis=1)`

In [78]: `df_master.head()`

Out[78]:

	MPG	Horsepower	Acceleration	log_Displacement	log_Weight	Cylinders_4	Cylinders_5	Cylinders_6
0	8.0	130.0	12.0	5.726848	8.161660	0	0	0
2	18.0	150.0	11.0	5.762051	8.142063	0	0	0
3	16.0	150.0	12.0	5.717028	8.141190	0	0	0
4	17.0	140.0	10.5	5.710427	8.145840	0	0	0
12	15.0	150.0	9.5	5.991465	8.232440	0	0	0

create x (with all independent variable) and y(with all target variable)

In [79]: `x=df_master.drop('MPG',axis=1)`

```
In [80]: y=df_master['MPG']
```

Random sampling :create training and test samples

```
In [81]: xtrain,xtest,ytrain,ytest=train_test_split(x,y,train_size=0.7,random_state=0)
```

feature selection

```
In [82]: #create key_features object to select the top k features  
# key_features = SelectKBest(score_func=f_regression,k='all')  
key_features=SelectKBest(score_func=f_regression,k=5)  
#to select 5 significant features  
  
#Fit the key_features to the training data and transform it  
xtrain_selected=key_features.fit_transform(xtrain,ytrain)  
  
#Get the indices of the selected features  
selected_indices=key_features.get_support(indices=True)  
  
#Get the names of the selected features  
selected_features=xtrain.columns[selected_indices]
```

```
In [83]: selected_features
```

```
Out[83]: Index(['Horsepower', 'log_Displacement', 'log_Weight', 'Cylinders_4',  
              'Cylinders_8'],  
              dtype='object')
```

Model 3: Build model using KBest selected features

```
In [84]: ##store KBest columns from xtrain to xtrain_kbest  
xtrain_kbest=xtrain[selected_features]
```

```
In [85]: #train your model  
linreg.fit(xtrain_kbest,ytrain)
```

```
Out[85]: LinearRegression()
```

```
In [86]: linreg.score(xtrain_kbest,ytrain)
```

```
Out[86]: 0.7118189417837262
```

```
In [88]: #store KBest columns from xtest to xtest_kbest
xtest_kbest=xtest[selected_features]
```

```
In [89]: pred_y=linreg.predict(xtest_kbest)
```

```
In [90]: linreg.score(xtest_kbest,ytest)
```

```
Out[90]: 0.7038559215939076
```

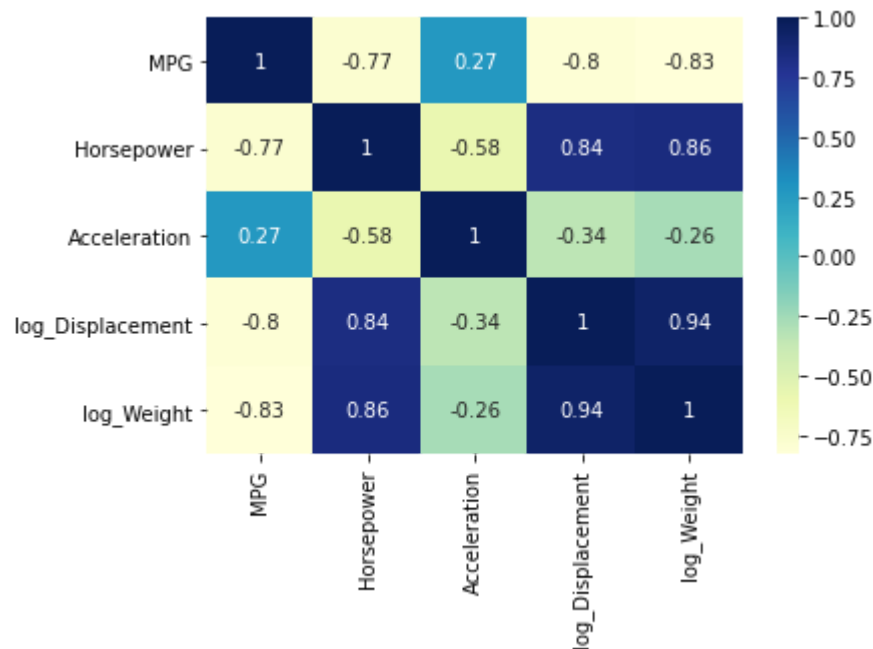
model 4

mutlicollinearity checking

```
In [91]: new_df_corr=df_numeric.corr()
```

```
In [92]: sns.heatmap(new_df_corr,cmap='YlGnBu',annot=True) #YlGnBu yellow green blue
```

```
Out[92]: <AxesSubplot:>
```



```
In [93]: #drop mutlicollunear variables from xtrain and xtest
xtrain=xtrain.drop(['log_Weight','log_Displacement'],axis=1)
xtest=xtest.drop(['log_Weight','log_Displacement'],axis=1)
```

feature selection¶

```
In [94]: #create key_features object to select the top k features  
# key_features = SelectKBest(score_func=f_regression,k='all')  
key_features=SelectKBest(score_func=f_regression,k=5)  
#to select 5 significant features  
  
#Fit the key_features to the training data and transform it  
xtrain_selected=key_features.fit_transform(xtrain,ytrain)  
  
#Get the indices of the selected features  
selected_indices=key_features.get_support(indices=True)  
  
#Get the names of the selected features  
selected_features=xtrain.columns[selected_indices]
```

```
In [95]: selected_features
```

```
Out[95]: Index(['Horsepower', 'Acceleration', 'Cylinders_4', 'Cylinders_8', 'Origin_3'], dtype='object')
```

Build model using KBest selected features

```
In [96]: ##store KBest columns from xtrain to xtrain_kbest  
xtrain_kbest=xtrain[selected_features]
```

```
In [97]: #train your model  
linreg.fit(xtrain_kbest,ytrain)
```

```
Out[97]: LinearRegression()
```

```
In [98]: linreg.score(xtrain_kbest,ytrain)
```

```
Out[98]: 0.7110316000206895
```

```
In [99]: #store KBest columns from xtest to xtest_kbest  
xtest_kbest=xtest[selected_features]
```

```
In [100]: pred_y=linreg.predict(xtest_kbest)
```

```
In [101]: linreg.score(xtest_kbest,ytest)
```

```
Out[101]: 0.703882872331421
```

```
In [ ]:
```

